

```
!pip install xgboost seaborn
```

```
Requirement already satisfied: xgboost in /usr/local/lib/python3.11/dist-packages (2.1.4)
Requirement already satisfied: seaborn in /usr/local/lib/python3.11/dist-packages (0.13.2)
Requirement already satisfied: numpy in /usr/local/lib/python3.11/dist-packages (from xgboost) (2.0.2)
Requirement already satisfied: nvidia-nccl-cu12 in /usr/local/lib/python3.11/dist-packages (from xgboost) (2.21.5)
Requirement already satisfied: scipy in /usr/local/lib/python3.11/dist-packages (from xgboost) (1.14.1)
Requirement already satisfied: pandas>=1.2 in /usr/local/lib/python3.11/dist-packages (from seaborn) (2.2.2)
Requirement already satisfied: matplotlib!=3.6.1,>=3.4 in /usr/local/lib/python3.11/dist-packages (from seaborn) (3.10.0)
Requirement already satisfied: contourpy>=1.0.1 in /usr/local/lib/python3.11/dist-packages (from matplotlib!=3.6.1,>=3.4->seaborn) (1.1.1)
Requirement already satisfied: cycler>=0.10 in /usr/local/lib/python3.11/dist-packages (from matplotlib!=3.6.1,>=3.4->seaborn) (0.12.1)
Requirement already satisfied: fonttools>=4.22.0 in /usr/local/lib/python3.11/dist-packages (from matplotlib!=3.6.1,>=3.4->seaborn) (4.53.0)
Requirement already satisfied: kiwisolver>=1.3.1 in /usr/local/lib/python3.11/dist-packages (from matplotlib!=3.6.1,>=3.4->seaborn) (1.4.6)
Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.11/dist-packages (from matplotlib!=3.6.1,>=3.4->seaborn) (24.1)
Requirement already satisfied: pillow>=8 in /usr/local/lib/python3.11/dist-packages (from matplotlib!=3.6.1,>=3.4->seaborn) (11.1.0)
Requirement already satisfied: pyparsing>=2.3.1 in /usr/local/lib/python3.11/dist-packages (from matplotlib!=3.6.1,>=3.4->seaborn) (3.1.4)
Requirement already satisfied: python-dateutil>=2.7 in /usr/local/lib/python3.11/dist-packages (from matplotlib!=3.6.1,>=3.4->seaborn) (2.9.0)
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.11/dist-packages (from pandas>=1.2->seaborn) (2025.2)
Requirement already satisfied: tzdata>=2022.7 in /usr/local/lib/python3.11/dist-packages (from pandas>=1.2->seaborn) (2025.2)
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.11/dist-packages (from python-dateutil>=2.7->matplotlib!=3.6.1,>=3.4->seaborn) (1.17.0)
```

```
# Import libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder, MinMaxScaler
from sklearn.metrics import mean_squared_error
from xgboost import XGBRegressor
```

```
import pandas as pd
path="/content/AmesHousing (1).csv"
data=pd.read_csv(path)
print(data)
```

```
Order      PID  MS SubClass  MS Zoning  Lot Frontage  Lot Area  Street
0          1  526301100      20      RL      141.0      31770  Pave
1          2  526350040      20      RH       80.0      11622  Pave
2          3  526351010      20      RL       81.0      14267  Pave
3          4  526353030      20      RL       93.0      11160  Pave
4          5  527105010      60      RL       74.0      13830  Pave
...      ...      ...      ...      ...      ...      ...
2925      2926  923275080      80      RL       37.0       7937  Pave
2926      2927  923276100      20      RL      NaN       8885  Pave
2927      2928  923400125      85      RL       62.0      10441  Pave
2928      2929  924100070      20      RL       77.0      10010  Pave
2929      2930  924151050      60      RL       74.0       9627  Pave
```

```
Alley  Lot Shape  Land Contour  ... Pool Area  Pool QC  Fence Misc Feature
0      NaN      IR1      Lvl ...      0      NaN      NaN      NaN
1      NaN      Reg      Lvl ...      0      NaN      MnPrv      NaN
2      NaN      IR1      Lvl ...      0      NaN      NaN      Gar2
3      NaN      Reg      Lvl ...      0      NaN      NaN      NaN
4      NaN      IR1      Lvl ...      0      NaN      MnPrv      NaN
...      ...      ...      ...      ...      ...      ...      ...
2925      NaN      IR1      Lvl ...      0      NaN      GdPrv      NaN
2926      NaN      IR1      Low ...      0      NaN      MnPrv      NaN
2927      NaN      Reg      Lvl ...      0      NaN      MnPrv      Shed
2928      NaN      Reg      Lvl ...      0      NaN      NaN      NaN
2929      NaN      Reg      Lvl ...      0      NaN      NaN      NaN
```

```
Misc Val Mo Sold Yr Sold Sale Type  Sale Condition  SalePrice
0          0      5      2010      WD      Normal      215000
1          0      6      2010      WD      Normal      105000
2      12500      6      2010      WD      Normal      172000
3          0      4      2010      WD      Normal      244000
4          0      3      2010      WD      Normal      189900
...      ...      ...      ...      ...      ...      ...
2925      0      3      2006      WD      Normal      142500
2926      0      6      2006      WD      Normal      131000
2927      700      7      2006      WD      Normal      132000
2928      0      4      2006      WD      Normal      170000
2929      0      11      2006      WD      Normal      188000
```

```
[2930 rows x 82 columns]
```

```
# Basic info
print("Dataset shape:", data.shape)
print("Columns:", data.columns.tolist())
```

```
Dataset shape: (2930, 82)
Columns: ['Order', 'PID', 'MS SubClass', 'MS Zoning', 'Lot Frontage', 'Lot Area', 'Street', 'Alley', 'Lot Shape', 'Land Contour', 'L
```

```
data.head()
```

	Order	PID	MS SubClass	MS Zoning	Lot Frontage	Lot Area	Street	Alley	Lot Shape	Land Contour	...	Pool Area	Pool QC	Fence	Misc Feature	Misc Val	Mo Sold
0	1	526301100	20	RL	141.0	31770	Pave	NaN	IR1	Lvl	...	0	NaN	NaN	NaN	0	5
1	2	526350040	20	RH	80.0	11622	Pave	NaN	Reg	Lvl	...	0	NaN	MnPrv	NaN	0	6
2	3	526351010	20	RL	81.0	14267	Pave	NaN	IR1	Lvl	...	0	NaN	NaN	Gar2	12500	6
3	4	526353030	20	RL	93.0	11160	Pave	NaN	Reg	Lvl	...	0	NaN	NaN	NaN	0	4
4	5	527105010	60	RL	74.0	13830	Pave	NaN	IR1	Lvl	...	0	NaN	MnPrv	NaN	0	3

5 rows × 82 columns

```
data.isnull().sum()
```

	0
Order	0
PID	0
MS SubClass	0
MS Zoning	0
Lot Frontage	490
...	...
Mo Sold	0
Yr Sold	0
Sale Type	0
Sale Condition	0
SalePrice	0

82 rows × 1 columns

```
data['Lot Frontage'].fillna(data['Lot Frontage'].mode()[0],inplace=True)
```

<ipython-input-13-951fcd9818e3>:1: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment. The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col]

```
data['Lot Frontage'].fillna(data['Lot Frontage'].mode()[0],inplace=True)
```

```
data.isnull().sum()
```

	0
Order	0
PID	0
MS SubClass	0
MS Zoning	0
Lot Frontage	0
...	...
Mo Sold	0
Yr Sold	0
Sale Type	0
Sale Condition	0
SalePrice	0


82 rows × 1 columns

```
# Encode categorical variables
le = LabelEncoder()
for col in data.columns:
    if data[col].dtype == "object":
        data[col] = le.fit_transform(data[col])

# Visualize normalized features (for learning)
scaler = MinMaxScaler()
scaled_data = pd.DataFrame(scaler.fit_transform(data), columns=data.columns)

plot_cols = ['Lot Area', 'Gr Liv Area', 'Total Bsmt SF', 'SalePrice']
plot_cols = [col for col in plot_cols if col in scaled_data.columns]

plt.figure(figsize=(10, 5))
sns.boxplot(data=scaled_data[plot_cols])
plt.title("Normalized Feature Distribution")
plt.grid()
plt.show()
```



```
# Split features and target
target_col = 'SalePrice'
X = data.drop(target_col, axis=1)
y = data[target_col]

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Train XGBoost Model with Tuned Parameters
xgb_model = XGBRegressor(
```

```

n_estimators=250,
learning_rate=0.07,
max_depth=6,
subsample=0.8,
colsample_bytree=0.8,
random_state=42
)

```

```
xgb_model.fit(X_train, y_train)
```

```

XGBRegressor
XGBRegressor(base_score=None, booster=None, callbacks=None,
              colsample_bylevel=None, colsample_bynode=None,
              colsample_bytree=0.8, device=None, early_stopping_rounds=None,
              enable_categorical=False, eval_metric=None, feature_types=None,
              gamma=None, grow_policy=None, importance_type=None,
              interaction_constraints=None, learning_rate=0.07, max_bin=None,
              max_cat_threshold=None, max_cat_to_onehot=None,
              max_delta_step=None, max_depth=6, max_leaves=None,
              min_child_weight=None, missing=nan, monotone_constraints=None,
              multi_strategy=None, n_estimators=250, n_jobs=None,
              num_parallel_tree=None, random_state=42, ...)

```

```
# Evaluate the model
```

```

y_pred = xgb_model.predict(X_test)
rmse = np.sqrt(mean_squared_error(y_test, y_pred))
print(f" RMSE of the XGBoost model: {rmse:.2f}")

```


```
RMSE of the XGBoost model: 23183.27
```

```
# Plot predicted vs actual prices
```

```

plt.figure(figsize=(8, 6))
plt.scatter(y_test, y_pred, alpha=0.6, color='teal')
plt.xlabel("Actual Sale Price")
plt.ylabel("Predicted Sale Price")
plt.title("Actual vs Predicted House Prices")
plt.grid(True)
plt.plot([y.min(), y.max()], [y.min(), y.max()], color='red', linestyle='--')
plt.show()

```



```
# : Feature Importance
```


```

plt.figure(figsize=(10, 6))
xgb_model.feature_importances_.argsort()
feat_imp = pd.Series(xgb_model.feature_importances_, index=X.columns)
feat_imp.nlargest(20).plot(kind='barh', color='skyblue')
plt.title("Top 20 Feature Importances - XGBoost")
plt.xlabel("Importance")
plt.grid()
plt.tight_layout()
plt.show()

```




Top 20 Feature Importances - XGBoost




```
# Residuals Plot
residuals = y_test - y_pred
plt.figure(figsize=(10, 5))
sns.histplot(residuals, bins=30, kde=True, color='coral')
plt.title("Distribution of Residuals (Actual - Predicted)")
plt.xlabel("Residual")
plt.grid()
plt.show()
```




Distribution of Residuals (Actual - Predicted)



```
# Prediction Error (Actual vs Predicted)
plt.figure(figsize=(8, 6))
sns.scatterplot(x=y_test, y=y_pred, color='navy')
plt.plot([y.min(), y.max()], [y.min(), y.max()], color='red', linestyle='--')
plt.xlabel("Actual Sale Price")
plt.ylabel("Predicted Sale Price")
plt.title("Prediction Error Plot")
plt.grid(True)
plt.tight_layout()
plt.show()
```



```
# Overlaid Histogram for Actual vs Predicted
plt.figure(figsize=(10, 5))
sns.histplot(y_test, label="Actual", kde=True, color='green', alpha=0.6)
sns.histplot(y_pred, label="Predicted", kde=True, color='orange', alpha=0.6)
plt.title("Actual vs Predicted Price Distribution")
plt.xlabel("Sale Price")
plt.legend()
plt.grid()
plt.show()
```



```
from sklearn.metrics import r2_score, mean_squared_error
import numpy as np
```

```
# Evaluate the model
r2 = r2_score(y_test, y_pred)
rmse = np.sqrt(mean_squared_error(y_test, y_pred))
```

```
print(" Final Model Evaluation:")
print(f" R2 Score (Accuracy): {r2:.4f}")
print(f" RMSE (Error in Price): {rmse:.2f}")
```



```
Final Model Evaluation:
R2 Score (Accuracy): 0.9330
RMSE (Error in Price): 23183.27
```